

HSLM-1.95M Architecture

Input: Token IDs (0-2047)

Embedding Layer

Weight: 2048×512 (ternary)
Positional: Learnable (512)

Transformer Stack (12 blocks)

Multi-Head Attention (8 heads)

Feed-Forward ($512 \rightarrow 1344 \rightarrow 512$) 90% sparse

Output Layer

Weight: $512 \rightarrow 2048$ (ternary)

Logits $\rightarrow$ Softmax